



CIS5200 Term Project Tutorial



Authors: [Yougender Chauhan](#), [Nagender Chauhan](#), [Suman Chauhan](#), [Vishwanth Reddy Sama](#).

Instructor: [Jongwook Woo](#)

Date: 12/16/2022

Lab Tutorial

Yougender Chauhan ([ychauha4@calstatela.edu](mailto:y Chauha4@calstatela.edu))

Nagender Chauhan (n Chauha5@calstatela.edu)

Suman Chauhan (s Chauha7@calstatela.edu)

Vishwanth Reddy Sama (vsama@calstatela.edu)

12/16/2022

Sentiment Analysis on Amazon Book Review Dataset

Objectives

In this hands-on lab, you will learn how to:

- Download Amazon book review data with -wget command.
- Upload the data into HDFS using -put command.
- Create Hive tables in HDFS using Hive/beeline QL
- Analysis on data using HIVE/beeline QL.
- Using a pre-determined Dictionary to do sentiment analysis on books reviews and their summary.

Platform Spec

- Cluster Version: Hadoop 3.1.2
- CPU Speed: 1995.309
- # of CPU cores: 4
- # of nodes: 5 (2 – Master and 3 – Worker)
- Total Memory Size: 390.71 GB

\$ hdfs version

\$ lscpu

\$ hdfs dfsadmin -report

```
-bash-4.2$ lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:             Little Endian
CPU(s):                 8
On-line CPU(s) list:   0-7
Thread(s) per core:    2
Core(s) per socket:    4
Socket(s):              1
NUMA node(s):          1
Vendor ID:              GenuineIntel
CPU family:             6
Model:                 85
Model name:             Intel(R) Xeon(R) Platinum 8167M CPU @ 2.00GHz
Stepping:              4
CPU MHz:               1995.309
BogoMIPS:               3990.61
Virtualization:        VT-x
Hypervisor vendor:     KVM
Virtualization type:   full
L1d cache:             32K
L1i cache:             32K
L2 cache:              4096K
L3 cache:              16384K
```

```
xgetbv1 xsaves arat umip pkru ospke md_clear arch_capabilities
-bash-4.2$ hdfs dfsadmin -report
Configured Capacity: 419520548352 (390.71 GB)
Present Capacity: 418101630960 (389.39 GB)
DFS Remaining: 147895687152 (137.74 GB)
DFS Used: 270205943808 (251.65 GB)
DFS Used%: 64.63%
Replicated Blocks:
  Under replicated blocks: 0
  Blocks with corrupt replicas: 0
  Missing blocks: 0
  Missing blocks (with replication factor 1): 0
  Low redundancy blocks with highest priority to recover: 0
  Pending deletion blocks: 0
Erasure Coded Block Groups:
  Low redundancy block groups: 0
  Block groups with corrupt internal blocks: 0
  Missing block groups: 0
  Low redundancy blocks with highest priority to recover: 0
  Pending deletion blocks: 0
report: Access denied for user schauha7. Superuser privilege is required
-bash-4.2$
```

```
Suman@DESKTOP-U61TL14 MINGW64 ~
$ ssh schauha7@144.24.14.145
schauha7@144.24.14.145's password:
Last login: Wed Dec  7 03:55:48 2022 from 047-041-218-030.res.spectrum.com
-bash-4.2$ hdfs version
Hadoop 3.1.2
Source code repository ssh://git@bitbucket.oc1.oraclecorp.com:7999/bdcs/apache_b
igtop.git -r 955ef423df4e67b7294f29b63c1e41eb6aec35e8
Compiled by root on 2022-10-26T22:15Z
Compiled with protoc 2.5.0
From source with checksum b367ca15864aef16725a3035859c9ece
This command was run using /usr/odh/1.1.2/hadoop/hadoop-common-3.1.2.jar
-bash-4.2$
```

Step 1: Downloading and uploading dataset into HDFS

Explain what this step is for. This step is to download the amazon book review dataset which is at my GitHub repository into the /tmp Linux server (since the file size is about 3gb) using wget command in gitbash CLI, then unzipping the dataset, uploading it into the HDFS.

1. Open Getbash CLI, connect to your Linux server using `$ ssh username@ipaddress`.

```
$ cd /tmp
$ wget https://github.com/Chauhan67/Amazonbookreviews/releases/download/v1/amazonbooks.zip
$ unzip amazonbooks.zip -d /tmp/books
$ cd /books
$ ls
$ hdfs dfs -mkdir books_data
$ hdfs dfs -mkdir Books_rating
$ hdfs dfs -put books_data.csv books_data
$ hdfs dfs -put Books_rating.csv Books_rating
$ cd
```

2. Run the above codes in this order to first get into /tmp directory to download .zip dataset, after which unzip it into /tmp/books sub-directory, create two directories in hdfs using `hdfs dfs -mkdir` command, upload into hdfs using `-put` command. Then use 'cd' in bash to go to your home directory i.e /home/<your-username>

```
$ wget -O dictionary.tsv https://github.com/dalgual/aidatasci/raw/master/data/bigdata/dictionary.tsv
$ hdfs dfs -mkdir dictionary
$ hdfs dfs -put dictionary.tsv dictionary
```

3. Now, download the dictionary in your Linux home directory, upload it into HDFS and view all files using `ls` and `cat` command.

Then you will see the following : `-bash-4.2$ hdfs dfs -ls`

```
$ hdfs dfs -ls
$ hdfs dfs -cat books_data/books_data.csv | head -n 10
$ hdfs dfs -cat Books_rating/Books_rating.csv | tail -n 10
```

```
Found 6 items
drwx---w- - ychauha4 hdfs      0 2022-12-07 05:05 .Trash
drwxr-xrwx - ychauha4 hdfs      0 2022-11-10 02:10 .hiveJars
drwxr-xr-x - ychauha4 hdfs      0 2022-12-03 00:07 Books_rating
drwxr-xr-x - ychauha4 hdfs      0 2022-12-03 00:07 books_data
drwxr-xr-x - ychauha4 hdfs      0 2022-12-03 01:08 dictionary
drwxr-xr-x - ychauha4 hdfs      0 2022-12-07 03:48 tmp
```

```
For-cat commands:-bash-4.2$ hdfs dfs -cat books_data/books_data.csv | head -n 10
```

Title,description,authors,image,previewLink,publisher,publishedDate,infoLink,categories,ratingsCount

```
Its      Only      Art      If      Its      well      Hung!,,['Julie Strain'],http://books.google.com/books/content?id=DykPAAAACAAJ&printsec=frontcover&img=1&zoom=1&source=gbs_api,http://books.google.nl/books?id=DykPAAAACAAJ&dq=Its+Only+Art+If+Its+well+Hung!&hl=&cd=1&source=gbs_api,,1996,http://books.google.nl/books?id=DykPAAAACAAJ&dq=Its+Only+Art+If+Its+well+Hung!&hl=&source=gbs_api,['Comics & Graphic Novels'],
```

```
-bash-4.2$ hdfs dfs -cat Books_rating/Books_rating.csv | tail -n 10
```

B000NSLVCU,The Idea of History,,AI1QNMVF2E3TN,"Robin George ""Master of Arts""",28/29,5.0,1057017600,R. G. Collingwood's Most Famous Book,"Highly Recommended.This book is one of the best books ever written on the Nature and Aims of History. This along with his "Principles of History" should give most readers all they need to know about the how and why of history.The book is extremely easy to read; harder to understand. Some criticisms of the book are not up to the mark, as for example complaints that Collingwood used Greek and Latin phrases in the book, and not everyone understands them.

Step 2: Data and sentiment analysis in Hive/beeline

Here, we will use hive or in our case beeline CLI to create tables and query data for analysis. We will be creating External Tables, Views, N-gram, Explode() etc for querying and analysis of the dataset.

```
-bash-4.2$ beeline
0: jdbc:hive2://bigdaiwn0.sub02180640120.traib> use <yourdatabase-name>;
```

You'll see the following:Driver: Hive JDBC (version 3.1.2)
Transaction isolation: TRANSACTION_REPEATABLE_READ
Beeline version 3.1.2 by Apache Hive
0: jdbc:hive2://bigdaiwn0.sub02180640120.traib>

1. Creating External Tables for books_data, Books_rating and dictionary by using out datasets. To do this use the following code:

```

CREATE EXTERNAL TABLE IF NOT EXISTS books_data (title STRING,
description STRING,
authors STRING,
image STRING,
preview STRING,
publisher STRING,
publish_date BIGINT,
info_link STRING,
categories STRING,
ratings_count float)
ROW FORMAT DELIMITED FIELDS TERMINATED BY ','
STORED AS TEXTFILE LOCATION '/user/ychauha4/books_data'
TBLPROPERTIES ('skip.header.line.count'='1');

```

NOTE – Replace ychauha4 with your username in the above code, after creating the table, to check we run the following select code:

```
select * from books_data limit 3;
```

You will see the following result with the column name – but not the header listed in the csv file:

```

+-----+-----+-----+-----+
+-----+-----+-----+-----+
+-----+-----+-----+-----+
+-----+-----+-----+-----+
|          books_data.title          |          books_data.description          |
|          books_data.authors        |          |          books_data.publisher          |
books_data.image                    |          |          books_data.publish_date        |
| books_data.publish_date            |          |          books_data.info_link          |
books_data.categories                |          |          |
+-----+-----+-----+-----+
+-----+-----+-----+-----+
+-----+-----+-----+-----+
+-----+-----+-----+-----+
| Its Only Art If Its well Hung!    |          |          |
| ['Julie Strain']                  |          |          |
http://books.google.com/books/content?id=source=gs_api |
http://books.google.nl/books?id=DykPAAACAAJ&dq=Its+Only+Art+If+Its+well+Hung!&hl=&cd=
1&source=gs_api |          | 1996
AAAACAAJ&dq=Its+Only+Art+If+Its+well+Hung!&hl=&source=gs_api | ['Comics & Graphic
Novels']          | NULL          |

```

```
| Dr. Seuss: American Icon | "Philip Nel takes a fascinating look into
the key aspects of Seuss's career - his poetry | politics
| a | marketing | and place in the popular
imagination." "Nel argues convincingly that Dr. Seuss is one of the most influential
poets in America. Hias changed language itself | giving us new
words like "nerd." And Seuss's famously loopy artistic style - what Nel terms an
"energetic cartoon surrealism" - has b |
| wonderful worship in Smaller Churches | "This resource includes twelve principles
in understanding small church worship | fifteen practices for planning worship with
fewer than 100 peop | ['David R. Ray']
http://books.google.com/books/content?id=2tsDAAAACAAJ&printsec=frontcover&img=1&zoom=1
&source=gbs_api | NULL | 2000
| NULL |
```

```
+-----+-----+-----+-----+
+-----+-----+-----+-----+
+-----+-----+-----+-----+
+-----+-----+-----+-----+
+-----+-----+-----+-----+
```

3 rows selected (0.333 seconds)

Create another external table for the books rating which has reviews for each book.

```
CREATE EXTERNAL TABLE IF NOT EXISTS books_rating (id BIGINT,
    title STRING,
    price FLOAT,
    user_id STRING,
    profile_name STRING,
    r_helpfulness STRING,
    r_score INT,
    r_time BIGINT,
    r_summary STRING,
    r_review STRING)
ROW FORMAT DELIMITED FIELDS TERMINATED BY ','
STORED AS TEXTFILE LOCATION '/user/ychauha4/Books_rating'
TBLPROPERTIES ('skip.header.line.count'='1');
```

NOTE – Replace ychauha4 with your username in the above code, after creating the table, to check we run the following select code:

```
select * from books_rating limit 3;
```

You will see the following result with the column name – but not the header listed in the csv file:

```
+-----+-----+-----+-----+
+-----+-----+-----+-----+
+-----+-----+-----+-----+
+-----+-----+-----+-----+
```

| books_rating.id | books_rating.title | books_rating.price | books_rating.user_id | books_rating.profile_name | books_rating.r_helpfulness | books_rating.r_score | books_rating.r_time | books_rating.r_summary | books_rating.r_review |
|-----------------|---|--------------------|----------------------|---------------------------|----------------------------|----------------------|---------------------|--|-----------------------|
| 1882931173 | Its Only Art If Its Well Hung! | | AVCGYZL8FQQT | "Jim of Oz ""jim-of-oz"" | 7/7 | 4 | | | |
| 940636800 | Nice collection of Julie Strain images | | | | | | | "This is only for Julie Strain fans. It's a collection of her photos -- about 80 pages worth with a nice section of paintings by Olivia.If you're looking for heavy literary content | |
| 826414346 | Dr. Seuss: American Icon | | A30TK6U7DNS82R | Kevin Killian | 10/10 | 5 | | | |
| 1095724800 | Really Enjoyed It | | | | | | | "I don't care much for Dr. Seuss but after reading Philip Nel's book I changed my mind--that's a good testimonial to the power of Rel's writing and thinking. Rel plays Dr. Seuss the ultimate compliment of treating him as a serious poet as well as one of the 20th century's most interesting visual artists | |
| 826414346 | Dr. Seuss: American Icon | | A3UH4UZ4RSV082 | John Granger | 10/11 | 5 | | | |
| 1078790400 | Essential for every personal and Public Library | | | | | | | "If people become the books they read and if ""the child is father to the man | |

3 rows selected (0.384 seconds)

Now Create a table for our dictionary which we will use for our sentiment analysis .

```
CREATE EXTERNAL TABLE if not exists dictionary (
    type string,
    length int,
    word string,
    pos string,
    stemmed string,
    polarity string )
ROW FORMAT DELIMITED
FIELDS TERMINATED BY '\t'
STORED AS TEXTFILE
LOCATION '/user/ychauha4/dictionary';
```

NOTE – Replace ychauha4 with your username in the above code, after creating the table, to check we run the following select code:

```
select * from dictionary limit 3;
```

```
+-----+-----+-----+-----+
| dictionary.type | dictionary.length | dictionary.word | dictionary.pos |
| dictionary.stemmed | dictionary.polarity |
+-----+-----+-----+-----+
| weaksubj      | 1 | | abandoned      | adj      | n
| negative      |   | |                 |          |
| weaksubj      | 1 | | abandonment    | noun     | n
| negative      |   | |                 |          |
| weaksubj      | 1 | | abandon        | verb     | y
| negative      |   | |                 |          |
+-----+-----+-----+-----+
3 rows selected (0.395 seconds)
```

Now, to see the number of records for books_data and books_rating tables we have created, use the following code:

```
select count(*) from books_data;

select count(*) from Books_rating;
```

```
0: jdbc:hive2://bigdaiwn0.sub02180640120.trai> select count(*) from books_data;
+-----+
| _c0 |
+-----+
| 212404 |
+-----+
1 row selected (13.111 seconds)
```

```
0: jdbc:hive2://bigdaiwn0.sub02180640120.trai> select count(*) from Books_rating;
+-----+
| _c0 |
+-----+
| 3000000 |
+-----+
1 row selected (12.699 seconds)
```

From this we know that we have 212404 different books and 3Million book reviews for the 212404 books.

2. Now, for the analysis we create a view to with title, average score and count of number of reviews for each book, we query this view to get most reviewed and least rated for more than 500 reviews then we use INSERTOVER DIRECTORY to save into our HDFS.

```
CREATE VIEW IF NOT EXISTS top10books AS
select title, avg(r_score) as avg_score, count(*) as count from books_rating group by title;
```

Next, to get top 10 rated books, their average rating and ordered by the count of reviews run this code:

```
INSERT OVERWRITE DIRECTORY '/user/ychauha4/tmp/top10'
ROW FORMAT DELIMITED FIELDS TERMINATED BY','
select * from top10books sort by count DESC limit 10;
```

NOTE - Change ychauha4 with your username.

You'll get the following output saved to /tmp/top10/000000_0;

```
-----
-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED
KILLED
-----
Map 1 ..... container  SUCCEEDED    1         1         0         0         0
0
Reducer 2 ..... container  SUCCEEDED   22        22         0         0         0
0
Reducer 3 ..... container  SUCCEEDED   11        11         0         0         0
0
Reducer 4 ..... container  SUCCEEDED    1         1         0         0         0
0
-----
VERTICES: 04/04 [=====>>] 100% ELAPSED TIME: 38.35 s
-----
No rows affected (44.282 seconds)
```

Now to get least 10 reviewed books with at least 500 reviews to it.

```
INSERT OVERWRITE DIRECTORY '/user/ychauha4/tmp/least10'
ROW FORMAT DELIMITED FIELDS TERMINATED BY','
select * from top10books where avg_score IS NOT NULL AND count >=500 sort by avg_score limit 10;
```

NOTE - Change ychauha4 with your username.

You'll get the following output saved to /tmp/least10/000000_0;

```
-----
-----
```

```

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED
KILLED
-----
Map 1 ..... container  SUCCEEDED    1         1         0         0         0
0
Reducer 2 ..... container  SUCCEEDED   22        22         0         0         0
0
Reducer 3 ..... container  SUCCEEDED    3         3         0         0         0
0
Reducer 4 ..... container  SUCCEEDED    1         1         0         0         0
0
-----
VERTICES: 04/04  [=====>>] 100%  ELAPSED TIME: 37.76 s
-----
No rows affected (43.226 seconds)

```

Now for top10categories by count run the following code:

```

INSERT OVERWRITE DIRECTORY '/user/ychauha4/tmp/top10categories'
ROW FORMAT DELIMITED FIELDS TERMINATED BY','
select categories, count(categories) as count from books_data where categories RLIKE
'(?<=\\').+?(?=\\)' group by categories sort by count DESC limit 10;

```

NOTE - Change ychauha4 with your username.

You'll get the following output saved to /tmp/top10categories/000000_0;

```

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED
KILLED
-----
Map 1 ..... container  SUCCEEDED    1         1         0         0         0
0
Reducer 2 ..... container  SUCCEEDED    6         6         0         0         0
0
Reducer 3 ..... container  SUCCEEDED    4         4         0         0         0
0
Reducer 4 ..... container  SUCCEEDED    1         1         0         0         0
0
-----
VERTICES: 04/04  [=====>>] 100%  ELAPSED TIME: 18.56 s
-----
No rows affected (23.865 seconds)

```

To check if the files have been downloaded to HDFS, in another gitbash window with your Linux system, run the below command:

```
$ hdfs dfs -ls tmp
$ hdfs dfs -cat tmp/top10/000000_0 |head -n 3
$ hdfs dfs -cat tmp/least10/000000_0 |head -n 3
$ hdfs dfs -cat tmp/top10categories/000000_0 |head -n 3
```

You should see the following outputs:

```
-bash-4.2$ hdfs dfs -ls tmp
Found 2 items
drwxr-xr-x - ychauha4 hdfs 0 2022-12-17 07:26 tmp/least10
drwxr-xr-x - ychauha4 hdfs 0 2022-12-17 07:19 tmp/top10
```

```
-bash-4.2$ hdfs dfs -cat tmp/top10/000000_0 |head -n 3
The Hobbit,4.655425755224971,22023
Pride and Prejudice,4.527145185990219,20371
Atlas Shrugged,4.026377349358715,12513
```

```
-bash-4.2$ hdfs dfs -cat tmp/least10/000000_0 |head -n 3
Red Rabbit,2.1573208722741435,643
Killing Time,2.3694029850746268,538
The Bear and the Dragon,2.4647627821280516,2179
```

```
-bash-4.2$ hdfs dfs -cat tmp/top10categories/000000_0 |head -n 3
['Fiction'],2123
['Religion'],1250
['History'],931
```

We will download this later with other files for visualization.

3. N-gram and context n-gram analysis of the most reviewed book i.e. 'The Hobbit' and least rated book 'Red Rabbit' first for 'The Hobbit' let us do a trigram. Run the below code:

```
SELECT EXPLODE(context_ngrams(sentences(LOWER(r_review)),
array(null, null, null), 10)) AS trigram FROM books_rating WHERE title = 'The Hobbit';
```

You should get the following output:

```
0: jdbc:hive2://bigdaiwn0.sub02180640120.trai> SELECT
EXPLODE(context_ngrams(sentences(LOWER(r_review)),
.> array(null, null, null), 10)) AS
trigram FROM books_rating WHERE title = 'The Hobbit';
```

| KILLED | VERTICES | MODE | STATUS | TOTAL | COMPLETED | RUNNING | PENDING | FAILED |
|---|----------|-----------|-----------|-------|-----------|---------|---------|--------|
| Map 1 | 0 | container | SUCCEEDED | 1 | 1 | 0 | 0 | 0 |
| Reducer 2 | 0 | container | SUCCEEDED | 1 | 1 | 0 | 0 | 0 |
| VERTICES: 02/02 [======>>] 100% ELAPSED TIME: 16.60 s | | | | | | | | |

```

+-----+
|               trigram               |
+-----+
| {"ngram":["lord","of","the"],"estfrequency":2670.0} |
| {"ngram":["of","the","rings"],"estfrequency":2457.0} |
| {"ngram":["the","lord","of"],"estfrequency":1871.0} |
| {"ngram":["the","hobbit","is"],"estfrequency":1693.0} |
| {"ngram":["read","this","book"],"estfrequency":1257.0} |
| {"ngram":["this","book","is"],"estfrequency":1020.0} |
| {"ngram":["read","the","hobbit"],"estfrequency":939.0} |
| {"ngram":["one","of","the"],"estfrequency":795.0} |
| {"ngram":["hobbit","is","a"],"estfrequency":734.0} |
| {"ngram":["this","is","a"],"estfrequency":685.0} |
+-----+
10 rows selected (22.103 seconds)

```

From this we can see a lot of tri-grams, to get further more context we now use context n-gram to find next 2 words after 'this' 'book' 'is' to get more insights by using the following code:

```
SELECT EXPLODE(context_ngrams(sentences(LOWER(r_review)),
array("this", "book", "is", null, null), 10)) AS fivegram FROM books_rating WHERE title = 'The Hobbit';
```

You should get the following output:

```
0: jdbc:hive2://bigdaiwn0.sub02180640120.trai> SELECT
EXPLODE(context_ngrams(sentences(LOWER(r_review)),
array("this", "book", "is", null,
null), 10)) AS fivegram FROM books_rating WHERE title = 'The Hobbit';
```

| KILLED | VERTICES | MODE | STATUS | TOTAL | COMPLETED | RUNNING | PENDING | FAILED |
|---|----------|-----------|-----------|-------|-----------|---------|---------|--------|
| Map 1 | 0 | container | SUCCEEDED | 1 | 1 | 0 | 0 | 0 |
| Reducer 2 | 0 | container | SUCCEEDED | 1 | 1 | 0 | 0 | 0 |
| VERTICES: 02/02 [======>>] 100% ELAPSED TIME: 18.45 s | | | | | | | | |

```

+-----+
|                    fivegram                    |
+-----+
| {"ngram":["about","a"],"estfrequency":50.0}    |
| {"ngram":["one","of"],"estfrequency":35.0}     |
| {"ngram":["a","must"],"estfrequency":24.0}     |
| {"ngram":["a","classic"],"estfrequency":20.0}  |
| {"ngram":["a","great"],"estfrequency":20.0}    |
| {"ngram":["the","best"],"estfrequency":20.0}   |
| {"ngram":["a","true"],"estfrequency":19.0}     |
| {"ngram":["a","wonderful"],"estfrequency":16.0}|
| {"ngram":["amazing","i"],"estfrequency":16.0}  |
| {"ngram":["a","timeless"],"estfrequency":15.0} |
+-----+
10 rows selected (19.078 seconds)

```

From this you can continue further with sixth gram and so on, with this fivegram we can say that most of them say it's a classic, great, best, wonderful etc... which is all positive.

Now for least 10 rated books i.e 'Red Rabbit'

```

SELECT EXPLODE(GRAMS(SENTENCES(LOWER(r_summary)), 2, 5))
AS bigrams
FROM books_rating
WHERE title = 'Red Rabbit';

```

You get the following output:

```

+-----+
|                    bigrams                    |
+-----+
| {"ngram":["red","rabbit"],"estfrequency":54.0} |
| {"ngram":["tom","clancy"],"estfrequency":28.0} |
| {"ngram":["jack","ryan"],"estfrequency":21.0}  |
| {"ngram":["of","the"],"estfrequency":11.0}    |
| {"ngram":["the","worst"],"estfrequency":9.0}   |
+-----+
5 rows selected (21.182 seconds)

```

From this we for further analysis this time let's see the summary and review where we have the words 'the worst' in it, run the following code:

```

SELECT r_summary, r_review
FROM books_rating
WHERE title = 'Red Rabbit'
AND r_summary LIKE '%the worst%' LIMIT 3;

```

You get the following output:

```

+-----+-----+
|          r_summary          |          r_review          |
+-----+-----+
| By far the worst Jack Ryan book | "I thought long and hard before |
| writing this review. I knew my complaints with Red Rabbit had much to do with my |
| knowledge of what Clancy was capable of when working at full capacity |
| One of the worst books I've ever struggled through. | "I like Tom Clancy. With the |
| exception of the new hardcover now out |
| Horrible.... the worst from Clancy EVER | "This book is horrible... even |
| the Bear and the Dragon was better than this (and I also disliked that one). Even if |
| we ignore Clancy's CONSTANT superficial right-wing rhetoric (not that is a bad |
| political stance |
+-----+-----+
3 rows selected (16.806 seconds)

```

We see that the same author has written another book ' the Bear and the Dragon' which is in fact 3rd least rated book from our analysis. This review backup out findings that 'red rabbit' is least rated.

4. Sentiment Analysis for book_rating by using dictionary to get the sentiments of review text and review summary columns. Use the following 3 Hive commands to create 3 views that will allow us to do that:

```
-- Create view l1 to compute sentiment for review text column
create view IF NOT EXISTS l1 as
    select id, user_id, words
    from books_rating
    lateral view explode(sentences(lower(r_review))) dummy as words;

-- Create view l2 from l1 to compute sentiment
create view IF NOT EXISTS l2 as
    select id, user_id, word
    from l1
    lateral view explode(words) dummy as word;

-- Create view l3 from l2 to compute sentiment
create view IF NOT EXISTS l3 as select
    id,
    user_id,
    l2.word,
    case d.polarity
    when 'negative' then -1
    when 'positive' then 1
    else 0 end as polarity
    from l2 left outer join dictionary d on l2.word = d.word;
```

Viewing the contents of l3 by using the following code:

```
select * from l3 limit 3;
```

You get the following output:

```
0: jdbc:hive2://bigdaiwn0.sub02180640120.trai> select * from l3 limit 3;
+-----+-----+-----+-----+
| 13.id | 13.user_id | 13.word | 13.polarity |
+-----+-----+-----+-----+
```

| | | | |
|------------|---------------|------|---|
| 1882931173 | AVCGYZL8FQQTD | this | 0 |
| 1882931173 | AVCGYZL8FQQTD | is | 0 |
| 1882931173 | AVCGYZL8FQQTD | only | 0 |

3 rows selected (10.451 seconds)

We create an external table called review_sentiment to sum all sentiments by the review id to get the sentiment of that review by using the below code:

```
create external table IF NOT EXISTS review_sentiment(id bigint,user_id bigint, r_sentiment int)
stored as orc;

INSERT OVERWRITE TABLE review_sentiment select

id,user_id,

case

when sum( polarity ) > 0 then '2'

when sum( polarity ) < 0 then '0'

else '1' end as sentiment

from l3 group by id,user_id;

select * from review_sentiment limit 3;
```

After successfully running these codes, when you run the select statement, you'll get the following output:

0: jdbc:hive2://bigdaiwn0.sub02180640120.trai> select * from review_sentiment limit 3;

| review_sentiment.id | review_sentiment.user_id | review_sentiment.r_sentiment |
|---------------------|--------------------------|------------------------------|
| NULL | NULL | 1 |
| NULL | NULL | 1 |
| NULL | NULL | 1 |

3 rows selected (0.495 seconds)

Do the same for summery text column by running the following codes:


```

-- Create view t1 to compute sentiment for summary text column
create view IF NOT EXISTS t1 as
select id, user_id, words
from books_rating
lateral view explode(sentences(lower(r_summary))) dummy as words;

-- Create view t2 from t1 to compute sentiment
create view IF NOT EXISTS t2 as
select id, user_id, word
from t1
lateral view explode(words) dummy as word;

-- Create view t3 from t2 to compute sentiment
create view IF NOT EXISTS t3 as select
id,
user_id,
t2.word,
case d.polarity
when 'negative' then -1
when 'positive' then 1
else 0 end as polarity
from t2 left outer join dictionary d on t2.word = d.word;

```

Viewing the contents of t3 by using the following code:

```
select * from t3 limit 3;
```

You get the following output:

```
0: jdbc:hive2://bigdaiwn0.sub02180640120.trai> select * from t3 limit 3;
```

| t3.id | t3.user_id | t3.word | t3.polarity |
|------------|--------------|------------|-------------|
| 1882931173 | AVCGYZL8FQQT | nice | 1 |
| 1882931173 | AVCGYZL8FQQT | collection | 0 |
| 1882931173 | AVCGYZL8FQQT | of | 0 |

3 rows selected (4.824 seconds)

We create an external table called `summary_sentiment` to sum all sentiments by the review id to get the sentiment of that review by using the below code:

```
create external table IF NOT EXISTS summary_sentiment(id bigint, user_id bigint, s_sentiment int)
stored as orc;

INSERT OVERWRITE TABLE summary_sentiment select
  id, user_id,
  case
  when sum( polarity ) > 0 then '2'
  when sum( polarity ) < 0 then '0'
  else '1' end as sentiment
from t3 group by id, user_id;

select * from summary_sentiment limit 3;
```

After successfully running these codes, when you run the select statement, you'll get the following output:

```
0: jdbc:hive2://bigdaiwn0.sub02180640120.trai>select * from summary_sentiment limit 3;
+-----+-----+-----+
| summary_sentiment.id | summary_sentiment.user_id | summary_sentiment.s_sentiment |
+-----+-----+-----+
| NULL                 | NULL                       | 1                             |
| NULL                 | NULL                       | 1                             |
| NULL                 | NULL                       | 1                             |
+-----+-----+-----+
3 rows selected (0.518 seconds)
```

Joining both review sentiment and summary sentiment by leftjoin over id and user_is using following code:

```
create view IF NOT EXISTS rs_sentiments as select
  r.id, r.user_id, r.r_sentiment, s.s_sentiment from review_sentiment r left outer join
  summary_sentiment s on r.id = s.id and r.user_id = s.user_id;

select * from rs_sentiments limit 3;
```

Let's see the output:

```
0: jdbc:hive2://bigdaiwn0.sub02180640120.trai> select * from rs_sentiments limit 3;
```

```

+-----+-----+-----+-----+
| rs_sentiments.id | rs_sentiments.user_id | rs_sentiments.r_sentiment | |
rs_sentiments.s_sentiment |
+-----+-----+-----+-----+
| NULL | NULL | 1 | NULL
| NULL | NULL | 1 | NULL
| NULL | NULL | 1 | NULL
+-----+-----+-----+-----+
3 rows selected (8.426 seconds)

```

Joining the review, summary sentiment and, books_rating into another table by creating an external book_senti table using the below code:

```

create external table IF NOT EXISTS books_senti(id bigint, title string, r_date timestamp, r_score
float, r_sentiment int, s_sentiment int) stored as orc;

```

```

INSERT OVERWRITE TABLE books_senti select
s.id,
r.title,
cast(from_unixtime(r_time) as timestamp) r_date,
r.r_score,
s.r_sentiment,
s.s_sentiment
FROM books_rating r LEFT OUTER JOIN rs_sentiments s on r.id = s.id;

```

```

Select * from books_senti limit 1;

```

The output will be:

```

0: jdbc:hive2://bigdaiwn0.sub02180640120.trai> Select * from books_senti limit 1;

```

```

+-----+-----+-----+-----+
| books_senti.id | books_senti.title | books_senti.r_date | |
books_senti.r_score | books_senti.r_sentiment | books_senti.s_sentiment |
+-----+-----+-----+-----+
| 1882931173 | Its Only Art If Its well Hung! | 1999-10-23 00:00:00.0 | 4.0
| 2 | NULL | |
+-----+-----+-----+-----+
1 row selected (0.388 seconds)

```

Run these two queries to see the deviation in sentiment tone for same book review and its summary:

```
select count(*) from books_senti;  
select count(*) from books_senti where review_sentiment != summary_sentiment;
```

```
0: jdbc:hive2://bigdaiwn0.sub02180640120.trai> select count(*) from books_senti;
```

```
+-----+  
|  _c0  |  
+-----+  
| 3252448 |  
+-----+
```

```
1 row selected (26.628 seconds)
```

```
0: jdbc:hive2://bigdaiwn0.sub02180640120.trai> select count(*) from books_senti where  
review_sentiment != summary_sentiment;
```

```
+-----+  
|  _c0  |  
+-----+  
| 298047 |  
+-----+
```

```
1 row selected (22.341 seconds)
```

From this you can see that in 3.2M reviews about 300k reviews have different sentiment for the same books review text and it's review summary column, this shows the limitations of our sentiment analysis where in it is difficult for machine to get 100% accurate sentiment. Major problem for this being double negation and sarcastic reviews which are sometimes even difficult for humans to understand.

5. In order to clean the data by eliminating non-alphabetical and non-numeric characters for visualization, we create one final table books_review and a view for top25 books, we join the table to only get the sentiments for all these top25books by using the below code:

```
create view top25 as
select * from top10books order by count DESC limit 25;

CREATE EXTERNAL TABLE IF NOT EXISTS books_review ( id bigint, title string, r_date timestamp,
r_score float, r_sentiment int, s_sentiment int)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ","
STORED AS TEXTFILE
LOCATION "/user/ychauha4/tmp/books_reviews";

INSERT OVERWRITE TABLE books_review
select b.id,
REGEXP_REPLACE(b.title, '[^a-zA-Z0-9]+' , ' ') title,
b.r_date, b.r_score,
b.review_sentiment, b.summary_sentiment from books_senti b LEFT SEMI JOIN top25 t on b.title =
t.title
where id IS NOT NULL
AND ( review_sentiment =0 or review_sentiment =1 or review_sentiment =2)
AND ( summary_sentiment =0 or summary_sentiment =1 or summary_sentiment =2);
```

You'll get the following output:

```
0: jdbc:hive2://bigdaiwn0.sub02180640120.trai> create view top25 as
DESC limit 25;
No rows affected (0.679 seconds)
```

```
0: jdbc:hive2://bigdaiwn0.sub02180640120.trai> INSERT OVERWRITE TABLE books_review
select b.id,
REGEXP_REPLACE(b.title, '[^a-zA-Z0-9]+' , ' ') title,
b.r_date, b.r_score,
b.review_sentiment, b.summary_sentiment
from books_senti b LEFT SEMI JOIN top25 t on b.title = t.title
where id IS NOT NULL
AND ( review_sentiment =0 or
review_sentiment =1 or review_sentiment =2)
AND ( summary_sentiment =0 or
summary_sentiment =1 or summary_sentiment =2);
```

| | VERTICES | MODE | STATUS | TOTAL | COMPLETED | RUNNING | PENDING | FAILED |
|--------|----------|-----------|-----------|-------|-----------|---------|---------|--------|
| KILLED | | | | | | | | |
| Map 3 | | container | SUCCEEDED | 1 | 1 | 0 | 0 | 0 |

```

Map 4 ..... container    SUCCEEDED      1          1          0          0          0
0
Map 5 ..... container    SUCCEEDED      1          1          0          0          0
0
Reducer 6 ..... container SUCCEEDED     21         21          0          0          0
0
Reducer 7 ..... container SUCCEEDED      1          1          0          0          0
0
Map 1 ..... container    SUCCEEDED      1          1          0          0          0
0
Reducer 2 ..... container SUCCEEDED      1          1          0          0          0
0

```

```

-----
VERTICES: 07/07 [======>>>] 100% ELAPSED TIME: 61.71 s
-----

```

```

-----
No rows affected (68.8 seconds)

```

```

0: jdbc:hive2://bigdaiwn0.sub02180640120.trai> select * from books_review limit 3;

```

```

+-----+-----+-----+-----+
| books_review.id | books_review.title | books_review.r_date | |
books_review.r_score | books_review.r_sentiment | books_review.s_sentiment |
+-----+-----+-----+-----+
| 736641238      | Little women      | 2004-02-05 00:00:00.0 | 5.0
| 2              | 2                 |                       |
| 736641238      | Little women      | 2004-01-14 00:00:00.0 | 4.0
| 2              | 2                 |                       |
| 736641238      | Little women      | 2003-12-10 00:00:00.0 | 4.0
| 2              | 2                 |                       |
+-----+-----+-----+-----+

```

```

3 rows selected (0.419 seconds)

```

Step 3: Downloading and Visualization

Now that we have done all our analysis and quering in Hive/beeline it's now time to download the files from HDFS to out linux to out local PC by using -get and SCP commands as foolows:

-get commands

```
$ hdfs dfs -get tmp/books_reviews/000000_0 books_reviews.csv
```

```
$ hdfs dfs -get tmp/top10/000000_0 top10reviewed.csv
```

```
$ hdfs dfs -get tmp/least10/000000_0 least10rated.csv
```

```
$ hdfs dfs -get tmp/top10categories/000000_0 top10categories.csv
```

Check in Linux system with ls command as follows:

```
-bash-4.2$ ls
books_reviews.csv  least10rated.csv  top10categories.csv
dictionary.tsv     __MACOSX          top10reviewed.csv
```

As you can see the files have been downloaded to Linux server, now let's open another GitBash terminal, don't ssh and using SCP we download this to our personal computer for visualization:

```
$ scp ychauha4@144.24.14.145:/home/ychauha4/books_reviews.csv .

$ scp ychauha4@144.24.14.145:/home/ychauha4/top10reviewed.csv .

$ scp ychauha4@144.24.14.145:/home/ychauha4/least10rated.csv .

$ scp ychauha4@144.24.14.145:/home/ychauha4/top10categories.csv .
```

Note – Replace ychauha4 with your username and IP address with the current IP. You will be asked password for every download.

```
$ scp ychauha4@144.24.14.145:/home/ychauha4/books_reviews.csv .
ychauha4@144.24.14.145's password:
books_reviews.csv                                100% 2095KB   1.8MB/s   00:01

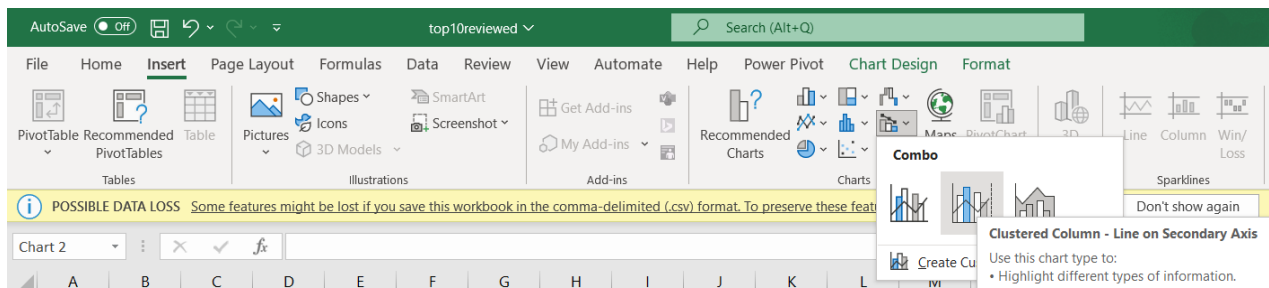
$ scp ychauha4@144.24.14.145:/home/ychauha4/top10reviewed.csv .
ychauha4@144.24.14.145's password:
top10reviewed.csv                                100% 416      5.2KB/s   00:00

$ scp ychauha4@144.24.14.145:/home/ychauha4/least10rated.csv .
ychauha4@144.24.14.145's password:
least10rated.csv                                100% 470      3.8KB/s   00:00

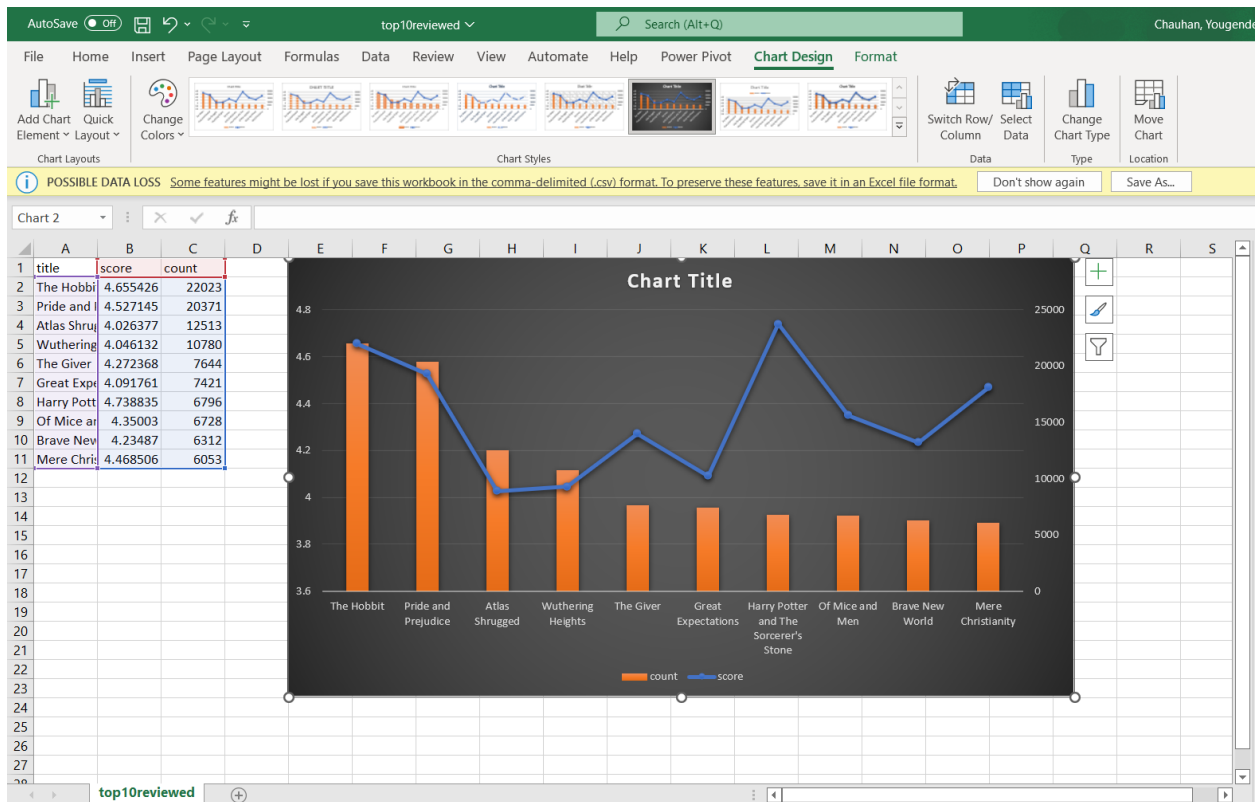
$ scp ychauha4@144.24.14.145:/home/ychauha4/top10categories.csv .
ychauha4@144.24.14.145's password:
top10categories.csv                              100% 232      5.2KB/s   00:00
```

Now to Visualize all these 4 documents in Excel and Tableau.

1. First let's start with our top10reviewed.csv file, open your Excel first then in Excel click on open file, navigate to your PC's file stored location in my case it is 'C:\Users\Yougander Chauhan'. Open the top10reviewed.csv.

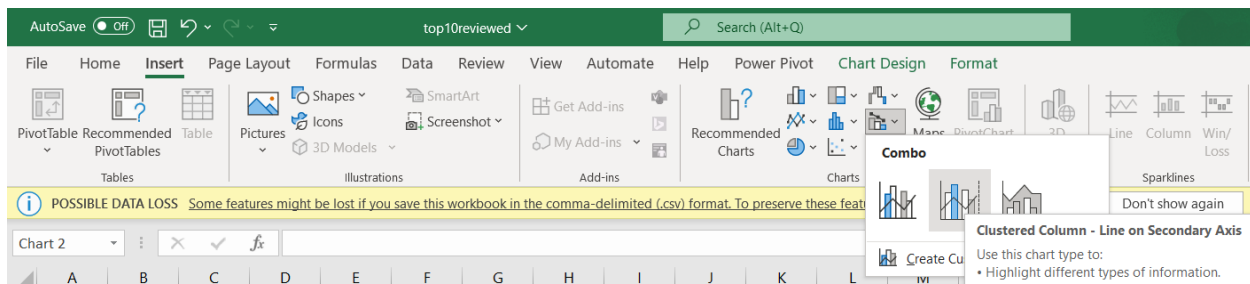


- After opening the excel, insert one top row, and add column names: title, score and count. Select all the filled rows and column, on top bar, click on 'insert' then on the right side, select Combo and then clustered column chart.

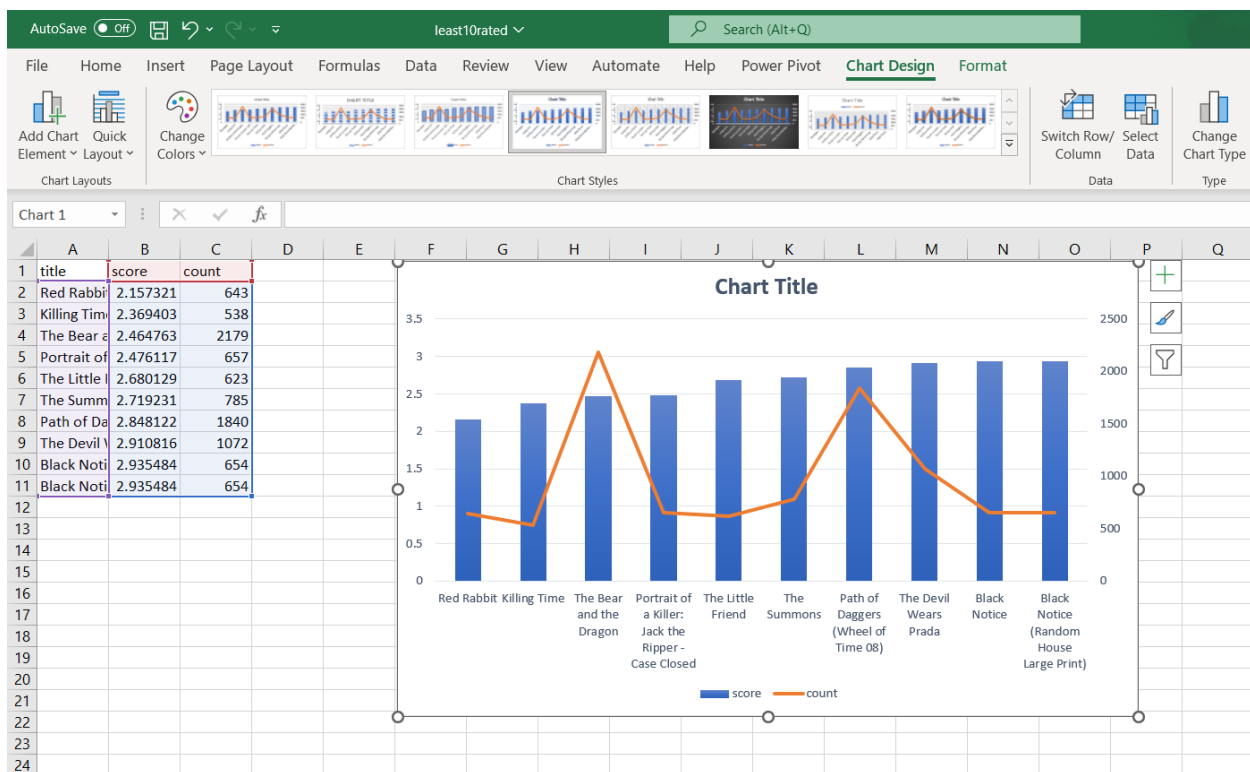


- This chart shows the Number of reviews for these 10 books and the line their respective average score.

2. Next, lets do the same for least10 rated books by following the above same steps. first then in Excel click on open file, navigate to your PC's file stored location in my case it is 'C:\Users\Yougender Chauhan'. Open the least10rated.csv.



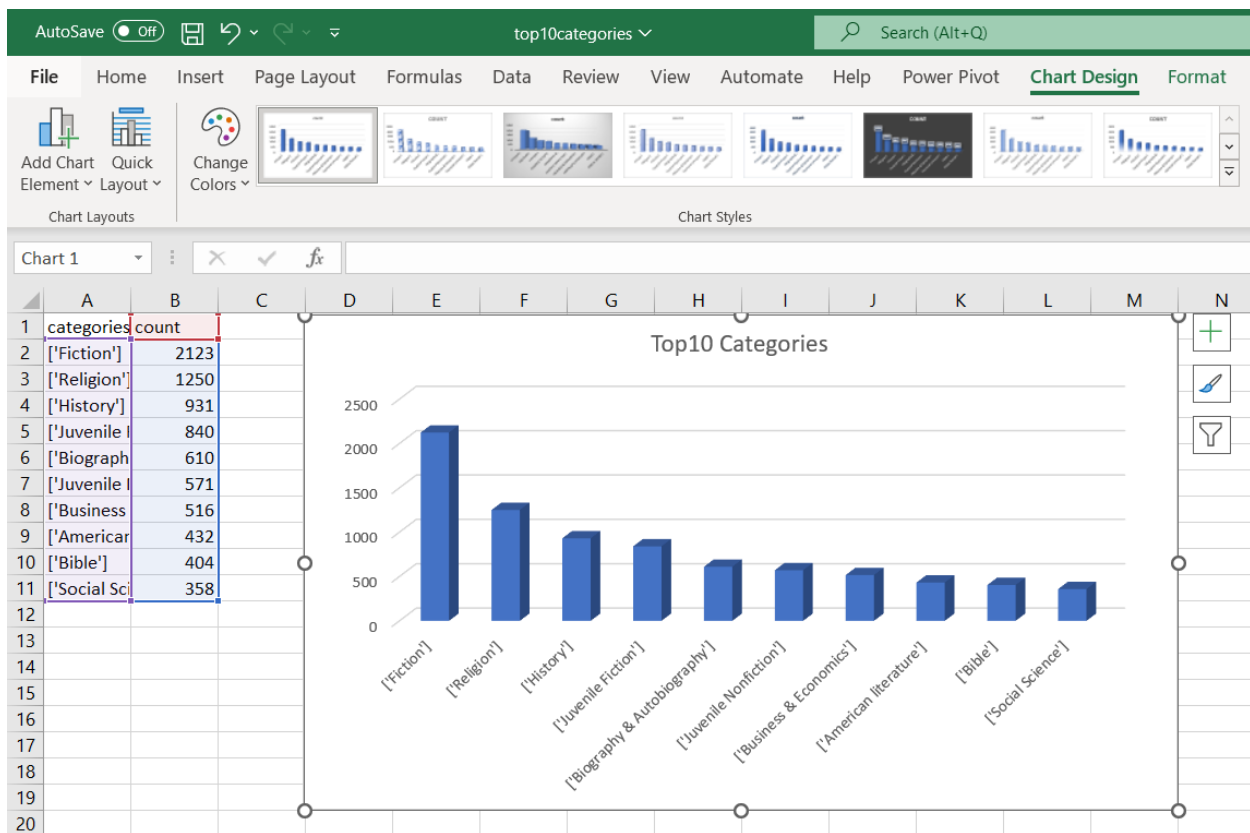
- After opening the excel, insert one top row, and add column names: title, score and count. Select all the filled rows and column, on top bar, click on 'insert' then on the right side, select Combo and then clustered column chart.



- This chart shows us the score and the line represent its respective number of reviews for that book.

3. Now for the top 10 categories of unique books present in our dataset. Open excel, click on open file then navigate to your PC's file stored location in my case it is 'C:\Users\Youngender Chauhan'. Open the top10categories.csv.

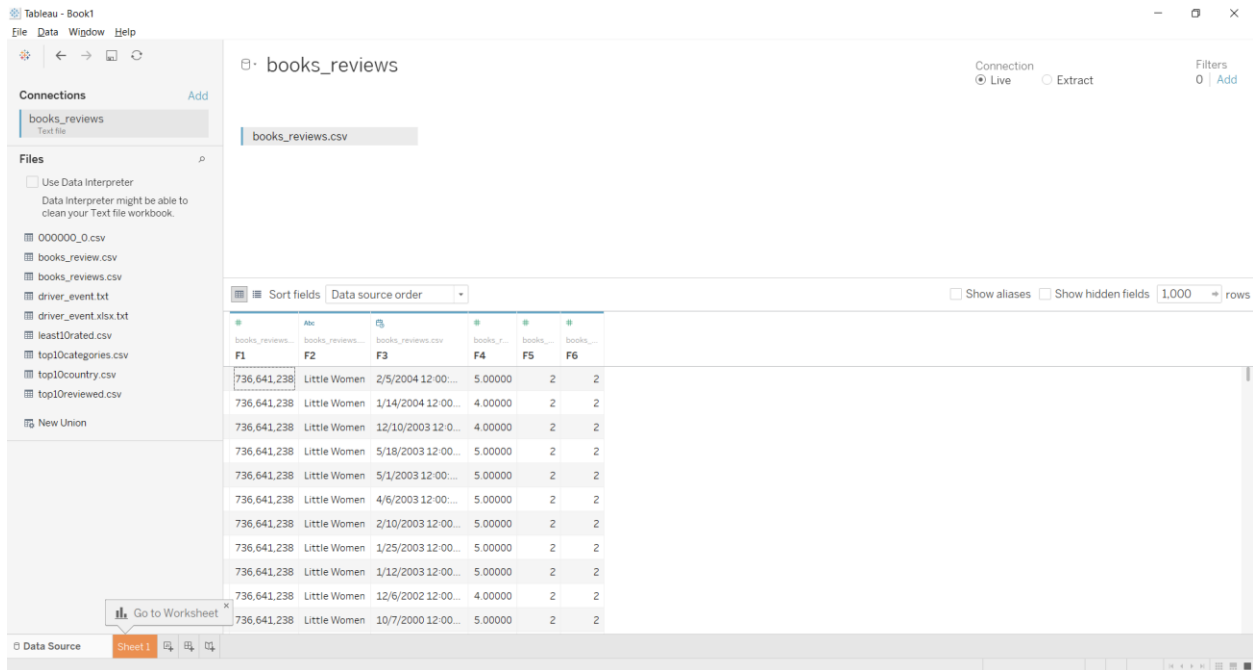
- After opening the excel, insert one top row, and add column names: categories and count. Select all the filled rows and column, on top bar, click on 'insert' then on the right side, select Bar Chart.



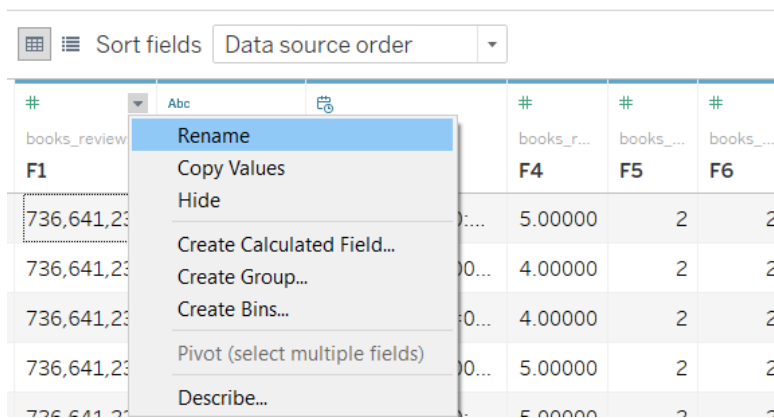
- This chart shows us the count of categories of each unique book.

4. Finally let's use tableau to visualize out sentiments that we have downloaded as 'books_reviews.csv'. to do so, first open then click on text file from there go to the path where the file is saved, in my case it is at 'C:\Users\Yougender Chauhan'. Open 'books_reviews.csv'.

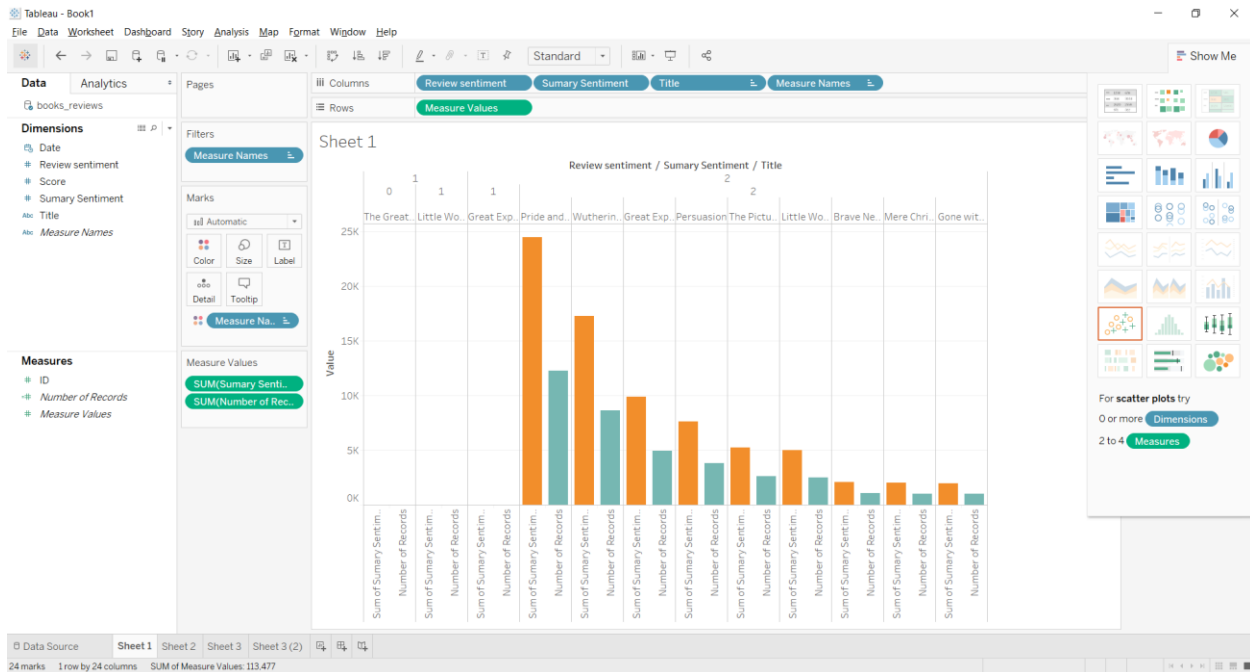
- When you load the data, it looks like this:



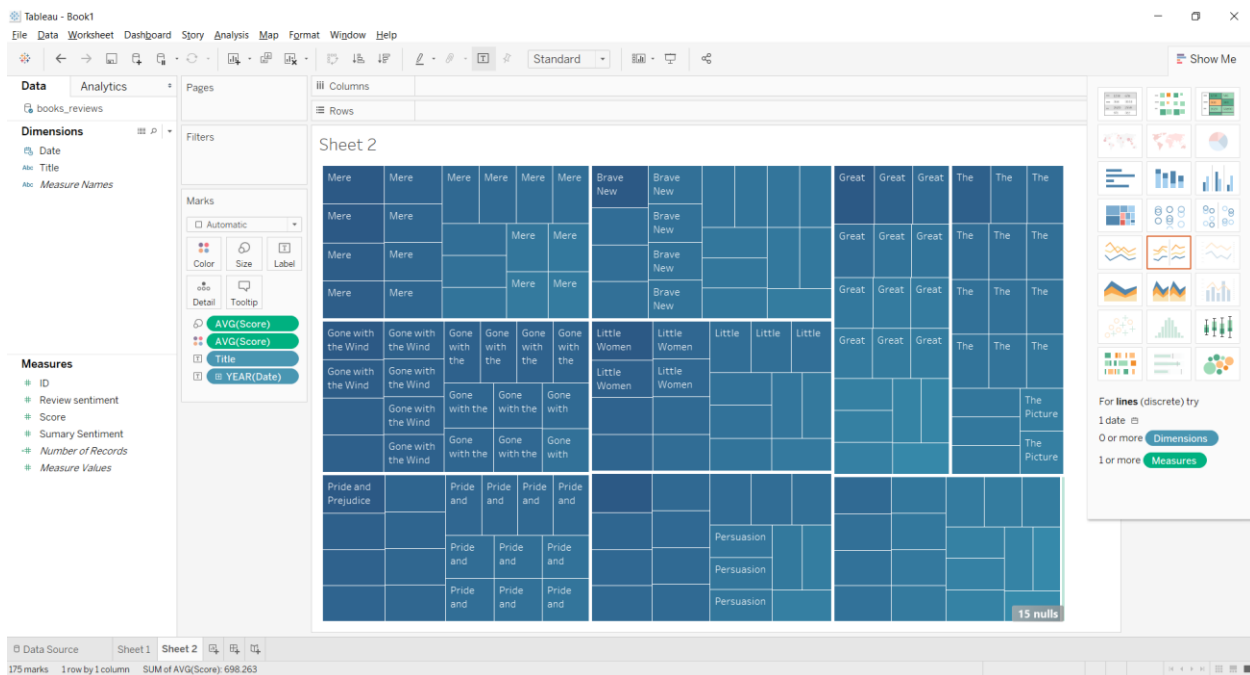
- Here click on dropdown on the column name F1 and click 'rename' and rename F1 to 'ID', F2 to 'Title', F3 to 'Date', F4 to 'Score', F5 to 'Review sentiment', F6 to 'Summary sentiment'. After which click on Sheet 1.



- Drag and drop 'Title' from Dimensions to Columns and 'Review Sentiment', 'Summary Sentiment' to Rows field. Move 'review sentiment' and 'summary sentiment' to dimensions and add to columns, then from Show Me drop down on right side, select side-by-side bar chart click on sort button from the top toolbar you get the following:

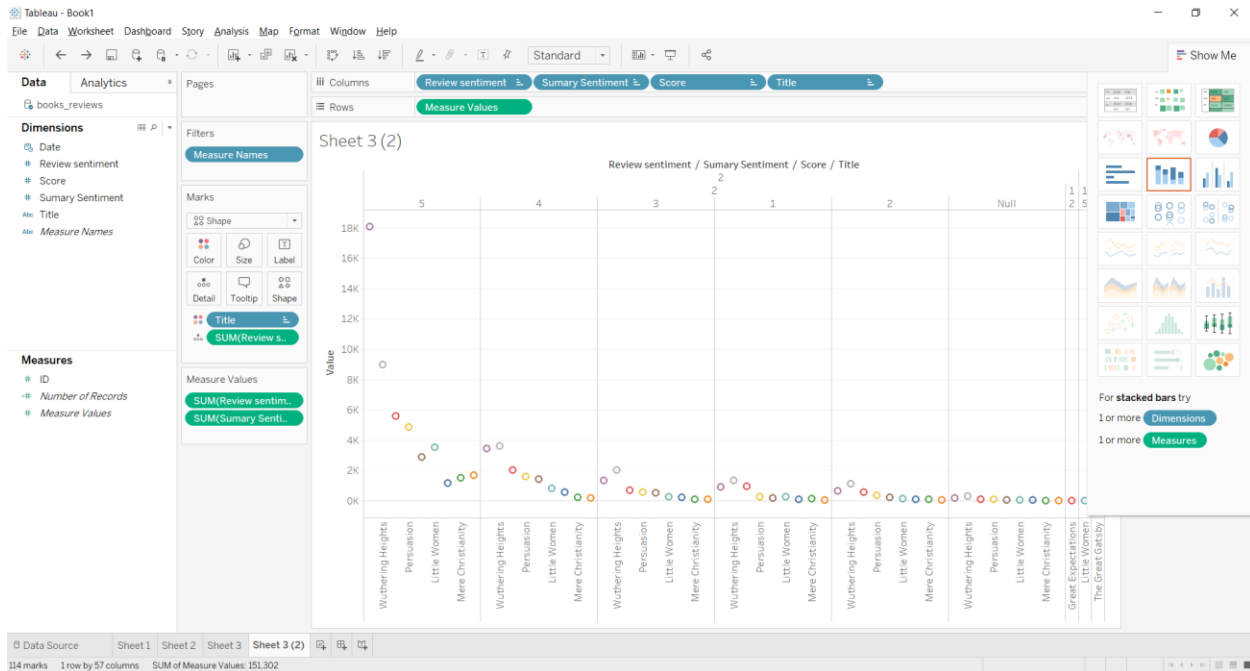


- Now, click on new worksheet beside 'sheet 1' on the bottom. On column add 'Title' and 'Date' on Rows add 'Score' then select treemaps from the show me dropdown on the right. Then in 'Marks' panel, sum(score) > measure(sum) > average. Do the same for color as well.



- You will see this treemaps which shows us the title and their average review score based on the year of review written.

- Next, click on new worksheet beside 'sheet 2' on the bottom. Move 'review sentiment' and 'summary sentiment' to dimensions, add 'score', 'title', 'review sentiment', 'summary sentiment'. Move 'review sentiment' and 'summary sentiment' to Measure and add it to rows. Select side-by-side circles on Show Me dropdown and sort by Title.



- This shows us the summary and review sentiments and the avg score side by side to see any deviations. Since we took a very small sample set for this example, we don't see the vast deviation here.

Summary

In this Tutorial you learned how to use Hadoop and Hive to do Sentiment analysis, Data analysis. We were also able to cover the limitation of using our current sentiment analysis model and using Excel and Tableau we visualized out analyzed data for better understanding.

References

1. URL of Data Source, https://www.kaggle.com/datasets/mohamedbakhmet/amazon-books-reviews?select=Books_rating.csv
2. URL of GitHub, <https://github.com/Chauhan67/Amazonbookreviews>

3. Lab tutorials used:

- labSentimentAnalysisTextNgrams
- labTwitterSentimentAnalysisLab
- labTableau_oracle_v2
- lab2HiveSensorDataAnalysisLab_aws